C35-NLP- Theory Session for

20/11/2022

Attention Models

Encoder-Decoder → Attention Models → Transformers → Hugging Face

Conceptual Session

Python Implementation

Introduction: Attention: most influential ideas in DL.

↳ NMT (Neural m/c Translation)

↳ Seq 2 Seq Models

↳ Encoder-Decoder architecture

English → Marathi

Encoder: process i/p seq & encode/compress/summarize info into CONTEXT/THOUGHT vector

Decoder: It is initialized with context vector, using which it starts generating transformed o/p.

I/p English: Rahul is a good boy.

Target Marathi: राहुल चांगला मुलगा आहे.

RNN: LSTM / GRU → complex sequence → large amt of data.

Appl.: Speech recognition, NLP, Time series forecasting, etc.

Seq2Seq

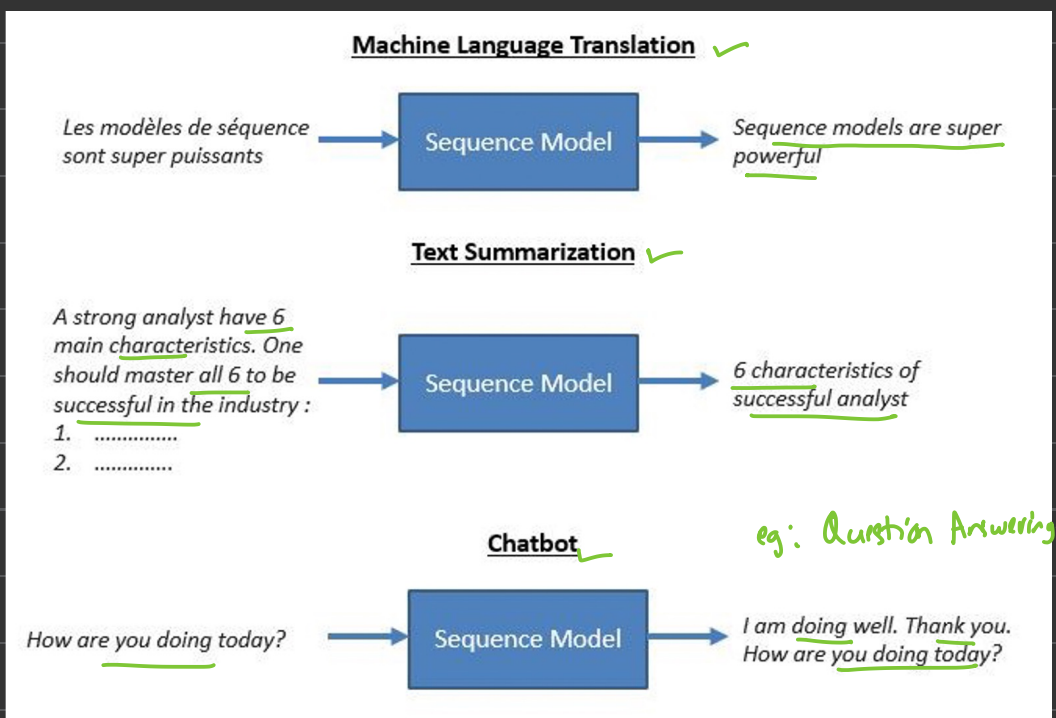

Prerequisites:

1. Fundamental concepts in ML & NN
2. Linear Algebra & Prob.
3. Working knowledge of LSTM.

Encoder - Decoder Architecture:

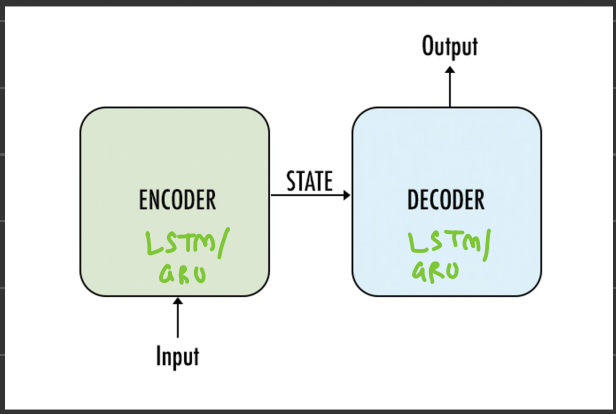

Encoder: read i/p seq.
↓
Summarize info as Internal state vectors (Hidden state & Cell state).
o/p of encoder are discarded & we only preserve the Internal State.

Decoder: I/p is final states of encoder
Using this, Decoder starts generating o/p seq.

Decoder Behaves differently in
- Training (uses Teacher-forcing)
- Inference (i/p to decoder at each time step is o/p from previous time step).

Encoder LSTM: 3 components: 1. X_i → i/p seq. at time step i
2. h_i & c_i → internal states
3. y_i → o/p seq. at time step i

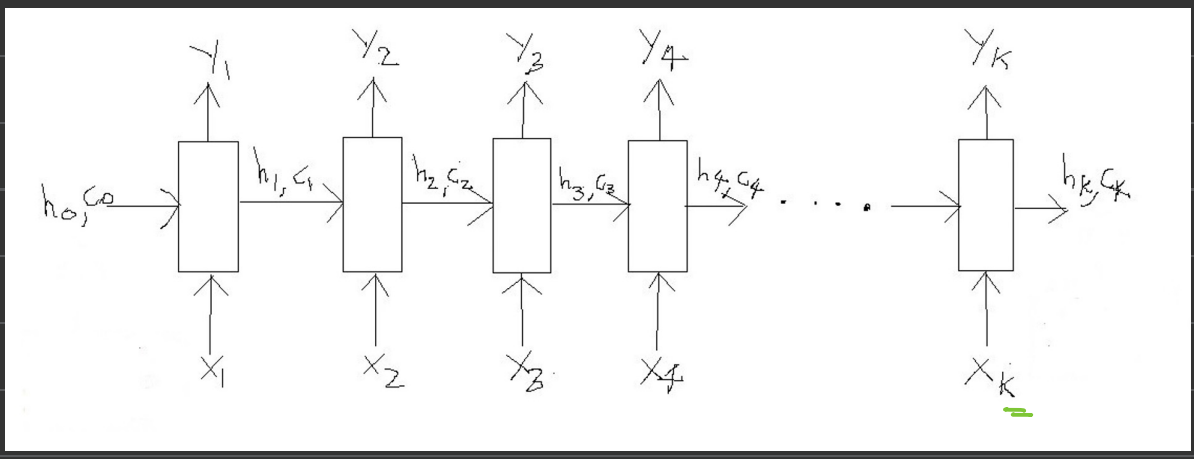

Word Level NMT:

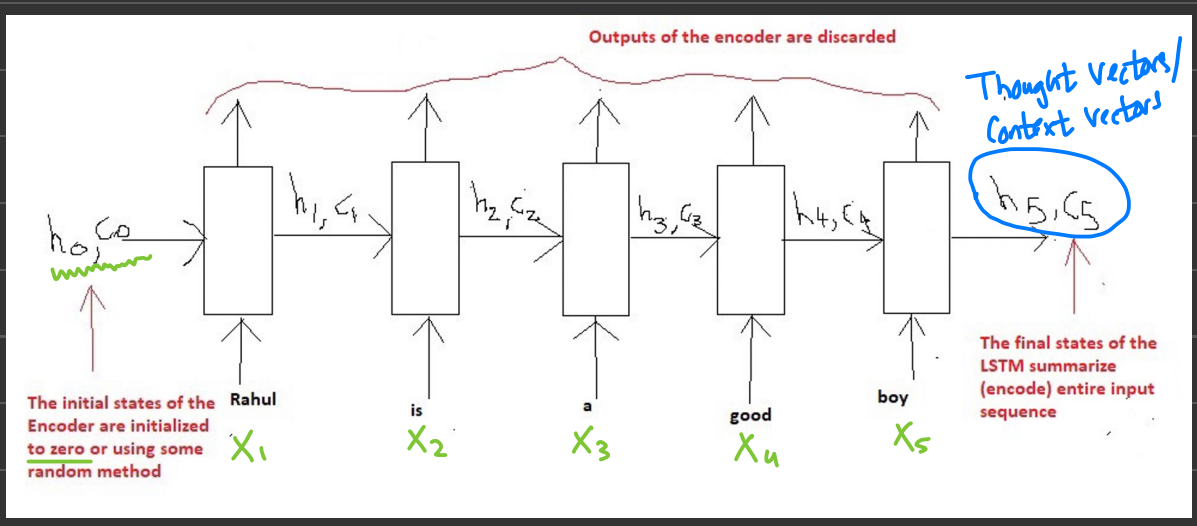

h_i & c_i: They remember what the LSTM has read (learned) till now

y_i : o/p (predictions) of LSTM at each time step)
y_i is prob. dist over entire vocabulary

Decoder LSTM – Training Mode.

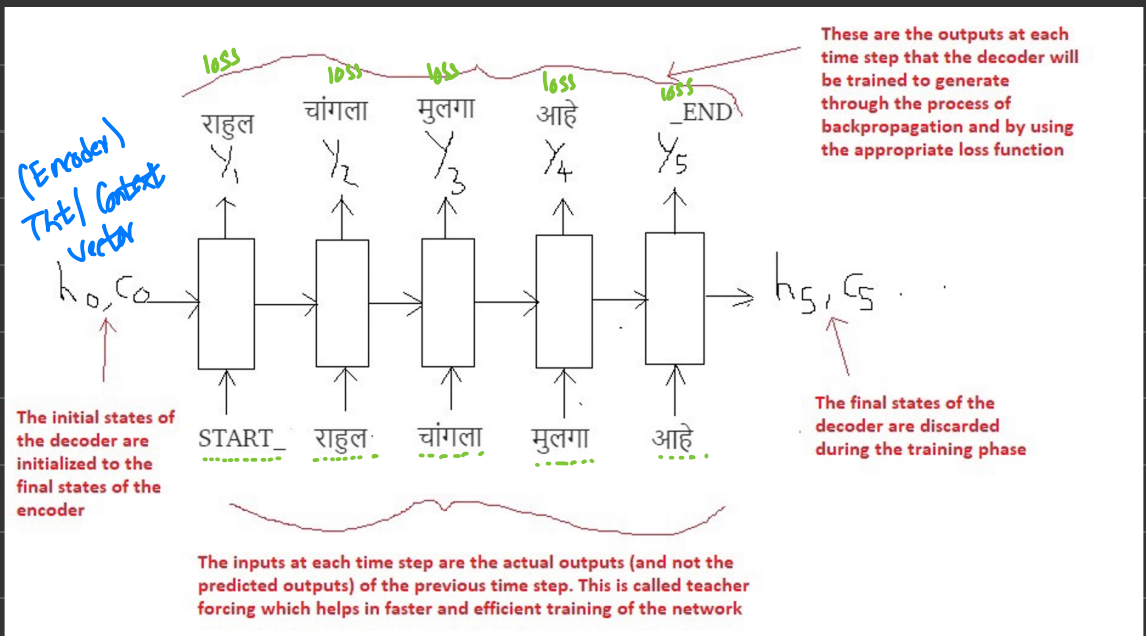

(TEACH)

Goal: Train the decoder to o/p राहुल चांगला मुलगा आहे.

"START_ राहुल चांगला मुलगा आहे _END"

Teacher Forcing: Here i/p at each time step is given as actual o/p (& not the predicted output) from the previous time step
– helps to train more fast.

https://machinelearningmastery.com/teacher-forcing-for-recurrent-neural-networks/

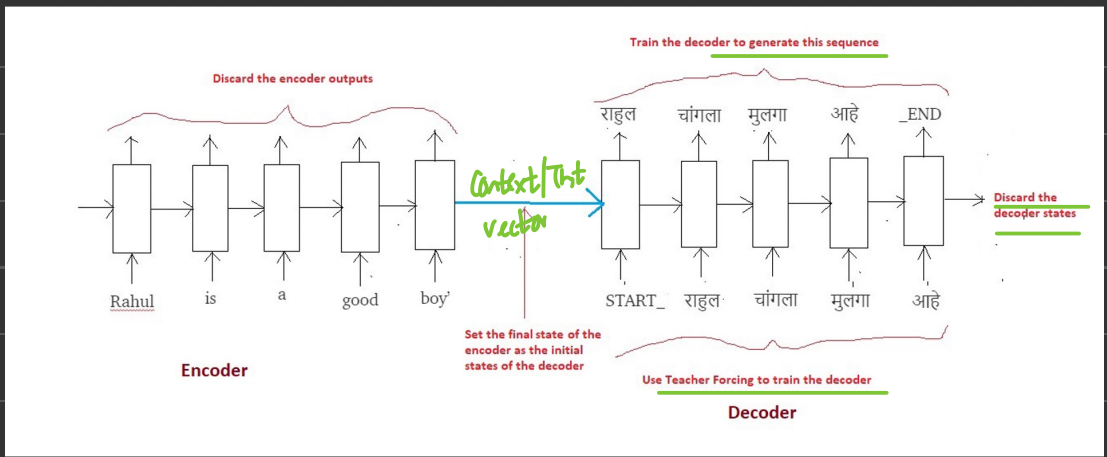

Decoder LSTM – Inference Mode:

1. Encode i/p seq. into Thought Vectors.

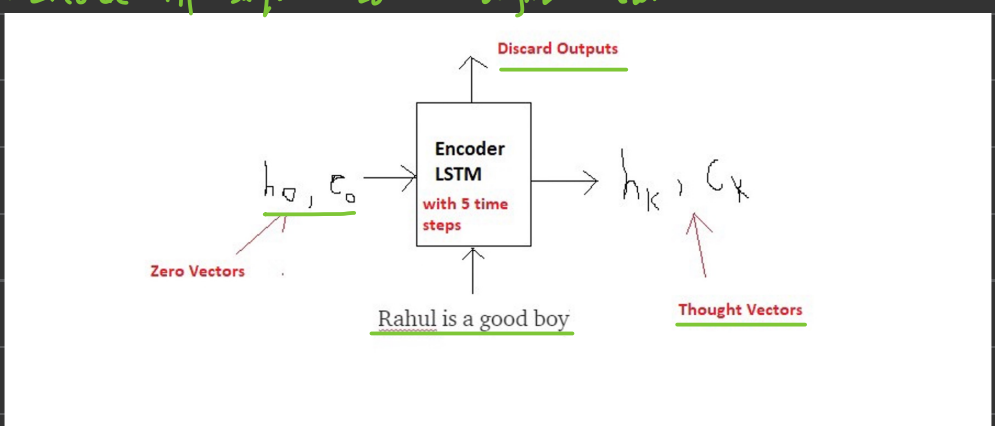

2. Start generating o/p seq. in a loop, word by word.

At $t=1$,

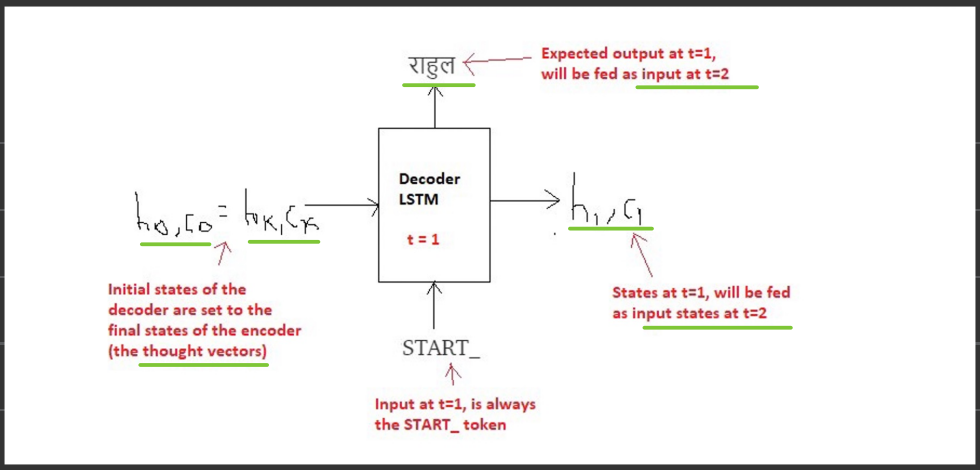
राहुल
Expected output at t=1, will be fed as input at t=2
Decoder LSTM
t = 1
Initial states of the decoder are set to the final states of the encoder (the thought vectors)
States at t=1, will be fed as input states at t=2
START_
Input at t=1, is always the START_ token

At t=2,

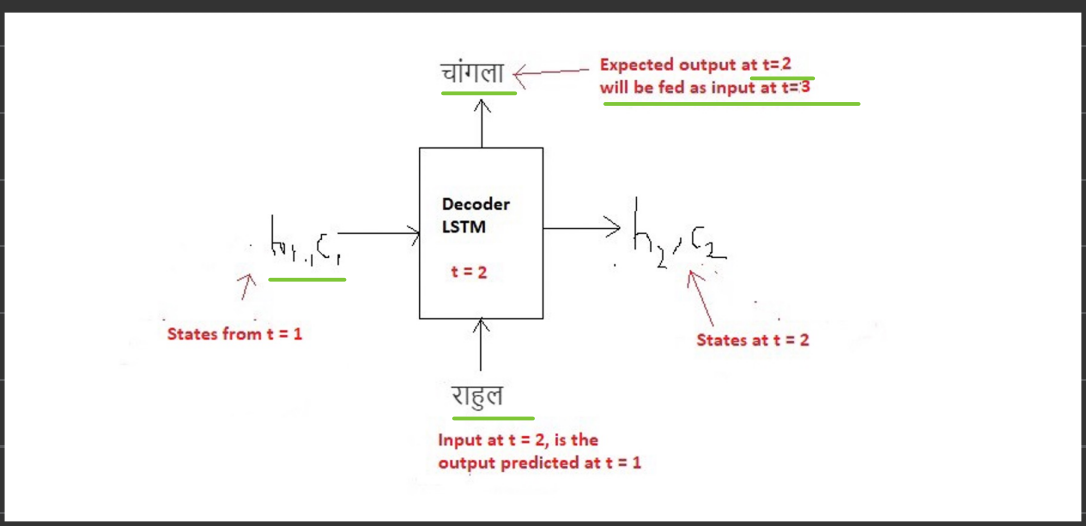
चांगला
Expected output at t=2 will be fed as input at t=3
Decoder LSTM
t = 2
States from t = 1
States at t = 2
राहुल
Input at t = 2, is the output predicted at t = 1

At t=3,

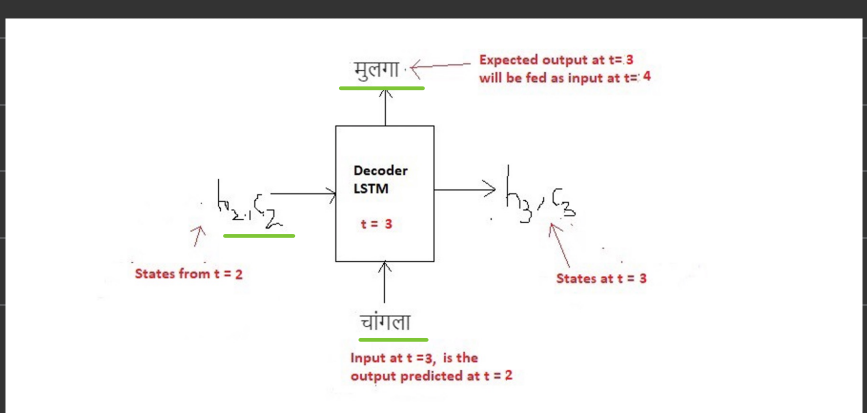
मुलगा
Expected output at t= 3 will be fed as input at t= 4
Decoder LSTM
t = 3
States from t = 2
States at t = 3
चांगला
Input at t =3, is the output predicted at t = 2

At t=4,

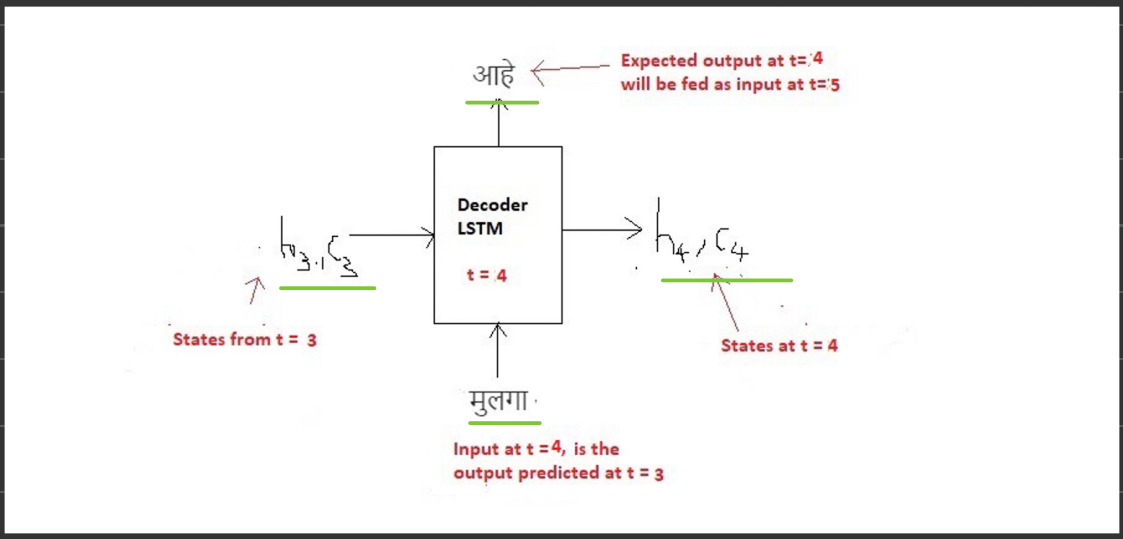
आहे
Expected output at t=4
will be fed as input at t=5
Decoder
LSTM
t = 4
h3, C3
h4, C4
States from t = 3
States at t = 4
मुलगा
Input at t =4, is the
output predicted at t = 3

At t=5,

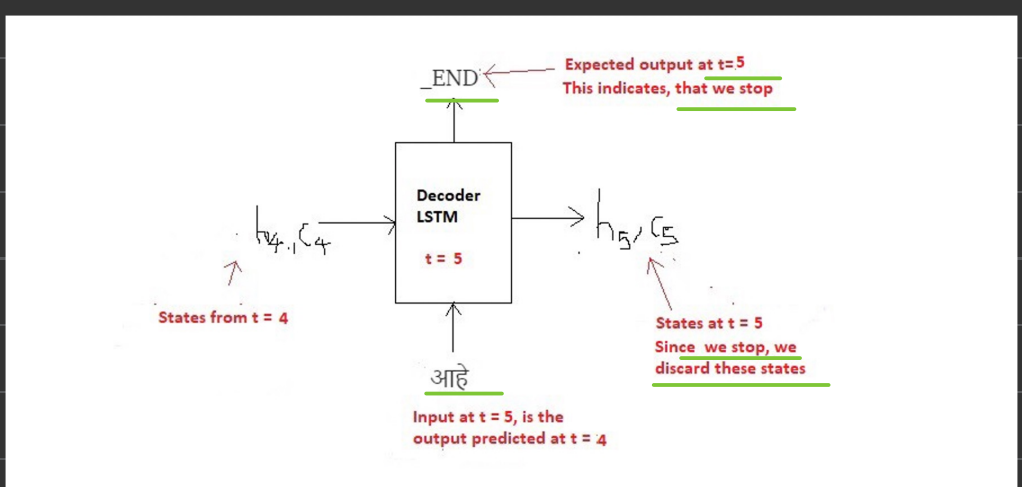
_END
Expected output at t=5
This indicates, that we stop
Decoder
LSTM
t = 5
h4, C4
h5, C5
States from t = 4
States at t = 5
Since we stop, we
discard these states
आहे
Input at t = 5, is the
output predicted at t = 4

Inference Algorithm:

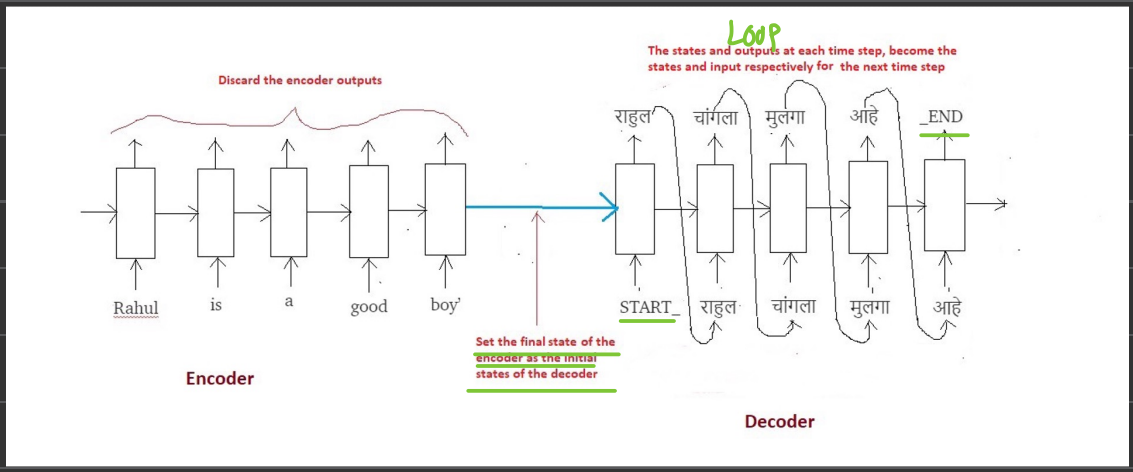

Why name Attention?

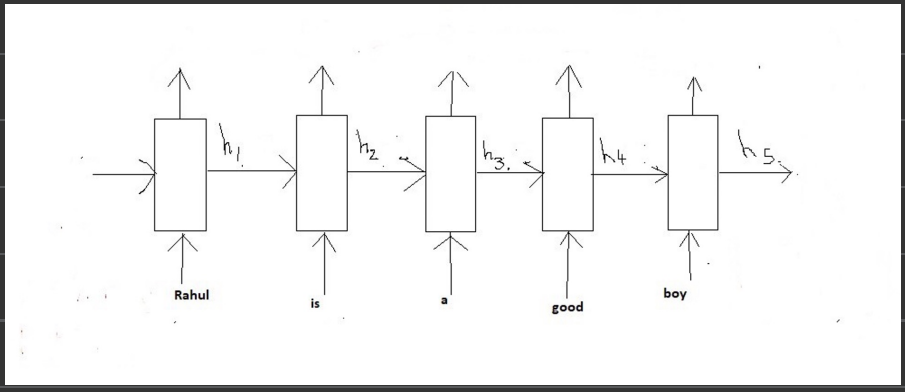

Encoder GRU:

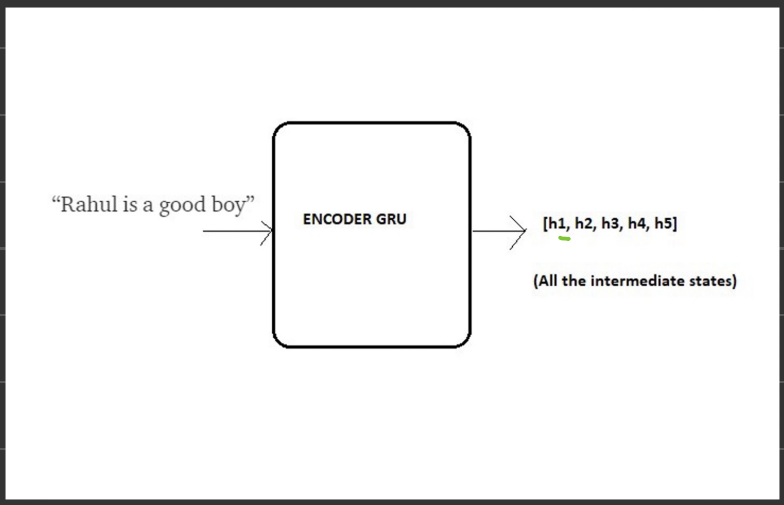

How does Attention work?

Decoding at time step 1.

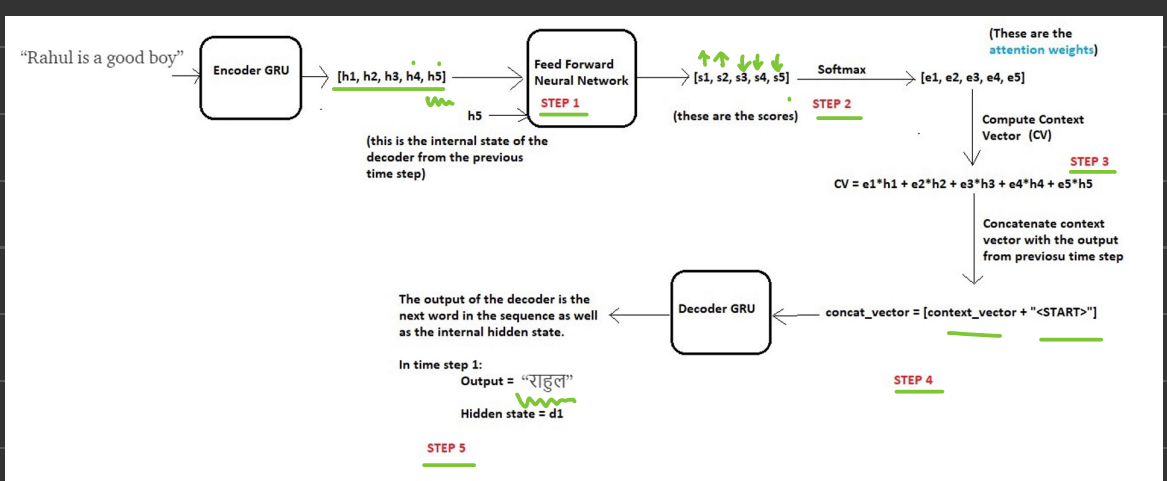

Compute a score "राहुल" Rahul

Compute attention weights : Apply softmax e1, e2, e3, e4, e5 [0-1]

$e_1 + e_2 + e_3 + e_4 + e_5 = 1$

$e_1 = 0.75,\ e_2 = 0.20,\ e_3 = 0.02,\ e_4 = 0.2,\ e_5 = 0.01$

Compute Context Vector $= e_1 \times h_1 + e_2 \times h_2 + e_3 \times h_3 + e_4 \times h_4 + e_5 \times h_5$

Concatenate context vector with o/p of prev. step